

Programming with fork/exec and pthreads

CIS 620 Advanced Operating Systems
Haodong Wang

1

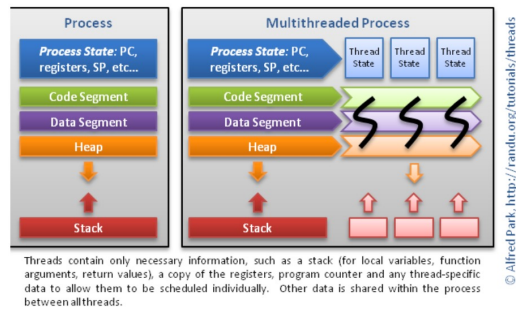
Processes

- Processes contain information about program resources and program execution state⁴, including:
 - Process ID, process group ID, user ID, and group ID
 - Environment, Working directory, Program instructions
 - Registers, Stack, Heap
 - Shared libraries, IPC tools (e.g., queues, pipes, semaphores, or shared memory).
- Process creation
 - Shell command: `./a.out`
 - Programming level: `fork()` system call

2

Threads

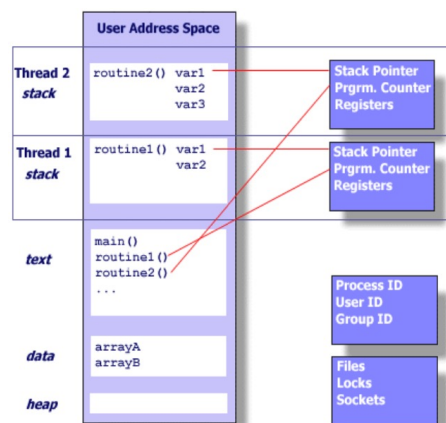
- Threads use, and exist within, the process resources
- Scheduled and run as independent entities
- Duplicate only the bare essential resources that enable them to exist as executable code



3

Threads

- A thread maintains its own:
 - Stack pointer
 - Registers
 - Scheduling properties
 - Signals (pending or blocked)
 - Thread specific data
- Multiple threads share the process resources
- A thread dies if the process dies
- “lightweight” for creating and terminating threads that for processes



4

What are pthreads?

- Threads used to implement parallelism in shared memory multiprocessor systems, such as SMPs
- Historically, hardware vendors have implemented their own proprietary versions of threads
- For Unix, a standardized C language threads programming interface has been specified by the IEEE Posix 1003.1c standard
 - pthreads
 - defines how threads should be created, managed, and destroyed

5

What are pthreads?

- Posix 1003.1c defines a thread interface
 - pthreads
 - defines how threads should be created, managed, and destroyed
- Unix provides a pthreads library
 - API to create and manage threads
 - you don't need to worry about the implementation details
 - this is a good thing

6

POSIX Thread API

- Commonly referred to as Pthreads, POSIX has emerged as the standard threads API, supported by most vendors.
 - Implemented with a `pthread.h` header/include file and a thread library
- Functionalities
 - Thread management, e.g., creation and joining
 - Thread synchronization primitives
 - Mutex
 - Conditional Variables
 - Reader/writer locks
 - Pthread barrier
 - Thread-specific data

7

PThread API

- `#include <pthread.h>`

Routine Prefix	Functional Group
<code>pthread_</code>	Threads themselves and miscellaneous subroutines
<code>pthread_attr_</code>	Thread attributes objects
<code>pthread_mutex_</code>	Mutexes
<code>pthread_mutexattr_</code>	Mutex attributes objects.
<code>pthread_cond_</code>	Condition variables
<code>pthread_condattr_</code>	Condition attributes objects
<code>pthread_key_</code>	Thread-specific data keys

- `gcc -lpthread`

8

Thread Creation

- Initially, main() program contains a single thread
 - All other threads must be explicitly created

```
int pthread_create(  
    pthread_t *thread,  
    const pthread_attr_t *attr,  
    void *(*start_routine) (void *),  
    void * arg);
```

- thread**: an opaque, unique identifier for the new thread returned by the subroutine
- attr**: an opaque attribute object that may be used to set thread attributes, e.g., NULL
- start_routine**: the C routine that the thread will execute once it is created
- arg**: a single argument that may be passed to start_routine. It must be passed by reference as a pointer cast of type void. NULL may be used if no argument is to be passed.

9

Example 1. pthread_create

```
#include <pthread.h>  
#define NUM_THREADS5  
  
void *PrintHello(void *thread_id) {  
    long tid = (long)thread_id;  
    printf("Hello World! It's me, thread #%ld!\n", tid);  
    pthread_exit(NULL);  
}  
  
int main(int argc, char *argv[]) {  
    pthread_t threads[NUM_THREADS];  
    long t;  
  
    for(t=0;t<NUM_THREADS;t++) {  
        printf("In main: creating thread %ld\n", t);  
        int rc = pthread_create(&threads[t], NULL, PrintHello, (void *)t);  
        if (rc) {  
            printf("ERROR: return code from pthread_create() is %d\n", rc);  
            exit(-1);  
        }  
    }  
    pthread_exit(NULL);  
}
```

One possible output:

```
In main: creating thread 0  
In main: creating thread 1  
In main: creating thread 2  
In main: creating thread 3  
Hello World! It's me, thread #0!  
In main: creating thread 4  
Hello World! It's me, thread #1!  
Hello World! It's me, thread #3!  
Hello World! It's me, thread #2!  
Hello World! It's me, thread #4!
```

10

Terminating Threads

- `pthread_exit` is used to explicitly exit a thread
 - Called after a thread has completed its work and is no longer required to exist
- If `main()` finishes before the threads it has created
 - If exits with `pthread_exit`, the other threads will continue to execute
 - Otherwise, they will be automatically terminated when `main()` finishes
- The programmer may optionally specify a termination status, which is stored as a void pointer for any thread that may join the calling thread
- Cleanup: the `pthread_exit()` routine does not close files
 - Any file opened inside the thread will remain open after the thread is terminated

11

Thread Attribute

```
int pthread_create(  
    pthread_t *thread,  
    const pthread_attr_t *attr,  
    void *(*start_routine) (void *),  
    void * arg);
```

- Attribute contains details about
 - Whether scheduling policy is inherited or explicit
 - Scheduling policy, scheduling priority
 - Stack size, stack guard region size
- `pthread_attr_init` and `pthread_attr_destroy` are used to initialize/destroy the thread attribute object
- Other routines are then used to query/set specific attributes in the thread attribute object

12

Passing Arguments to Threads

- The `pthread_create()` routine permits the programmer to pass **one** argument to the thread start routine
- For cases where multiple arguments must be passed:
 - Create a structure which contains all of the arguments
 - Then pass a pointer to the object of that structure in the `pthread_create()` routine
 - All arguments must be passed by reference and cast to `(void *)`
- Make sure that all passed data is thread safe
 - It cannot be changed by other threads
 - It can be changed by a determinant way

13

Example 2: Argument Passing

```
#include <pthread.h>
#define NUM_THREADS 8

struct thread_data {
    int thread_id;
    char *message;
};

struct thread_data thread_data_array[NUM_THREADS];

void *PrintHello(void *threadarg) {
    int taskid;
    char *hello_msg;

    sleep(1);
    struct thread_data *my_data = (struct thread_data *) threadarg;
    taskid = my_data->thread_id;
    hello_msg = my_data->message;
    printf("Thread %d: %s\n", taskid, hello_msg);
    pthread_exit(NULL);
}
```

14

Example 2: Argument Passing

```
int main(int argc, char *argv[]) {
    pthread_t threads[NUM_THREADS];
    int t;
    char *messages[NUM_THREADS];
    messages[0] = "English: Hello World!";
    messages[1] = "French: Bonjour, le monde!";
    messages[2] = "Spanish: Hola al mundo";
    messages[3] = "Klingon: Nuq neH!";
    messages[4] = "German: Guten Tag, Welt!";
    messages[5] = "Russian: Zdravstvuyte, mir!";
    messages[6] = "Japan: Sekai e konnichiwa!";
    messages[7] = "Latin: Orbis, te saluto!";

    for(t=0;t<NUM_THREADS;t++) {
        struct thread_data * thread_arg = &thread_data_array[t];
        thread_arg->thread_id = t;
        thread_arg->message = messages[t];
        pthread_create(&threads[t], NULL, PrintHello, (void *) thread_arg);
    }
    pthread_exit(NULL);
}
```


```
Thread 3: Klingon: Nuq neH!
Thread 0: English: Hello World!
Thread 1: French: Bonjour, le monde!
Thread 2: Spanish: Hola al mundo
Thread 5: Russian: Zdravstvuyte, mir!
Thread 4: German: Guten Tag, Welt!
Thread 6: Japan: Sekai e konnichiwa!
Thread 7: Latin: Orbis, te saluto!
```

15

Wait for Thread Termination

- Suspend execution of calling thread until thread terminates


```
#include <pthread.h>
int pthread_join(
    pthread_t thread,
    void **value_ptr);
```
- Thread: the joining thread
- Value_ptr: ptr to location for return code a terminating thread passes to pthread_exit



16

Thread Joining Example

```
#include <stdio.h>
#include <pthread.h>

void printMsg(char* msg) {
    printf("%s\n", msg);
}

int main(int argc, char** argv) {
    pthread_t thrdID;


    printf("creating a new thread\n");
    pthread_create(&thrdID, NULL, (void*)printMsg, argv[1]);
    printf("created thread %d\n", thrdID);
    pthread_join(thrdID, NULL);

    return 0;
}
```

17

Shared Memory and Threads

- Threads have access to the same global, shared memory
- Threads also have their own private data
- Programmers are responsible for synchronizing access globally shared data



18

Thread Caveat

- Shared State!
 - Accidental changes to global variables can be fatal
 - Changes made by one thread to shared resources (such as closing a file) will be seen by all other threads
 - Reading and writing to the same memory locations is possible
 - Therefore requires explicit synchronization by the programmer
- Many library functions are not thread-safe
 - Library functions that return pointers to static internal memory. E.g., `gethostbyname()`
- Lack of robustness
 - Crash in one thread will crash the entire process

19

Why Pthreads (not processes)?

- The primary motivation
 - To realize potential program performance gains
- Compared to the cost of creating and managing a process
 - A thread can be created with much less OS overhead
- Managing thread requires fewer system resources
- All threads within a process share the same address space
- Inter-thread communication is more efficient and, in many cases, easier than IPC

20

Pthread_create vs. Fork

- Timing results for the fork() system call and pthread_create() API
 - Timings reflect 50k process/thread creations
 - Units are in seconds
 - No optimization flags

Platform	fork()			pthread_create()		
	real	user	sys	real	user	sys
AMD 2.4 GHz Opteron (8cpus/node)	41.07	60.08	9.01	0.66	0.19	0.43
IBM 1.9 GHz POWER5 p5-575 (8cpus/node)	64.24	30.78	27.68	1.75	0.69	1.10
IBM 1.5 GHz POWER4 (8cpus/node)	104.05	48.64	47.21	2.01	1.00	1.52
INTEL 2.4 GHz Xeon (2 cpus/node)	54.95	1.54	20.78	1.64	0.67	0.90
INTEL 1.4 GHz Itanium2 (4 cpus/node)	54.54	1.07	22.22	2.03	1.26	0.67

21

Synchronizing Threads

- Three basic synchronization primitives
 1. mutex locks
 2. condition variables
 3. semaphores
- Mutexes and condition variables will handle most of the cases you need in this class
 - but feel free to use semaphores if you like

22

Mutex Locks

- A Mutex lock is created like a normal variable
 - `pthread_mutex_t mutex;`
- Mutexes must be initialized before being used
 - a mutex can only be initialized once
 - prototype:
 - `int pthread_mutex_init(pthread_mutex_t *mp, const pthread_mutexattr_t *mattr);`
 - *mp*: a pointer to the mutex lock to be initialized
 - *mattr*: attributes of the mutex – usually NULL

23

Locking a Mutex

- To insure mutual exclusion to a critical section, a thread should lock a mutex
 - when locking function is called, it does not return until the current thread owns the lock
 - if the mutex is already locked, calling thread blocks
 - if multiple threads try to gain lock at the same time, the return order is based on priority of the threads
 - higher priorities return first
 - no guarantees about ordering between same priority threads
 - prototype:
 - `int pthread_mutex_lock(pthread_mutex_t *mp);`
 - *mp*: mutex to lock

24

Unlocking a Mutex

- When a thread is finished within the critical section, it needs to release the mutex
 - calling the unlock function releases the lock
 - then, any threads waiting for the lock compete to get it
 - very important to remember to release mutex
 - prototype:
 - `int pthread_mutex_unlock(pthread_mutex_t *mp);`
 - *mp*: mutex to unlock

25

Example

```
#include <stdio.h>
#include <pthread.h>

#define MAX_SIZE 5
pthread_mutex_t bufLock;
int count;

void producer(char* buf) {
    for(;;) {
        while(count == MAX_SIZE);
        pthread_mutex_lock(&bufLock);
        buf[count] = getChar();
        count++;
        pthread_mutex_unlock(&bufLock);
    }
}

void consumer(char* buf) {
    for(;;) {
        while(count == 0);
        pthread_mutex_lock(&bufLock);
        useChar(buf[count-1]);
        count--;
        pthread_mutex_unlock(&bufLock);
    }
}

int main() {
    char buffer[MAX_SIZE];
    pthread_t p;
    count = 0;
    pthread_mutex_init(&bufLock);
    pthread_create(&p, NULL, (void*)producer, &buffer);
    consume(&buffer);
    return 0;
}
```

26

Condition Variables (CV)

- Notice in the previous example a *spin-lock* was used wait for a condition to be true
 - the buffer to be full or empty
 - spin-locks require CPU time to run
 - waste of cycles
- Condition variables allow a thread to block until a specific condition becomes true
 - recall that a blocked process cannot be run
 - doesn't waste CPU cycles
 - blocked thread goes to wait queue for condition
- When the condition becomes true, some other thread signals the blocked thread(s)

27

Condition Variables (CV)

- A CV is created like a normal variable
 - *pthread_cond_t condition;*
- CVs must be initialized before being used
 - a CV can only be initialized once
 - prototype:
 - `int pthread_cond_init(pthread_cond_t *cv, const pthread_condattr_t *cattr);`
 - *cv*: a pointer to the condition variable to be initialized
 - *cattr*: attributes of the condition variable – usually NULL

28

Blocking on CV

- A wait call is used to block a thread on a CV
 - puts the thread on a wait queue until it gets signaled that the condition is true
 - even after signal, condition may still not be true!
 - blocked thread does not compete for CPU
 - the wait call should occur under the protection of a mutex
 - this mutex is automatically released by the wait call
 - the mutex is automatically reclaimed on return from wait call
- prototype:
 - `int pthread_cond_wait(pthread_cond_t *cv, pthread_mutex_t *mutex);`
 - *cv*: condition variable to block on
 - *mutex*: the mutex to release while waiting

29

Signaling a Condition

- A signal call is used to “wake up” a single thread waiting on a condition
 - multiple threads may be waiting and there is no guarantee as to which one wakes up first
 - thread to wake up does not actually wake until the lock indicated by the wait call becomes available
 - condition thread was waiting for may not be true when the thread actually gets to run again
 - should always do a wait call inside of a while loop
 - if no waiters on a condition, signaling has no effect
 - prototype:
 - `int pthread_cond_signal(pthread_cond_t *cv);`
 - *cv*: condition variable to signal on

30

```

#include <stdio.h>
#include <pthread.h>

#define MAX_SIZE 5
pthread_mutex_t lock;
pthread_cond_t notFull, notEmpty;
int count;

void producer(char* buf) {
    for(;;) {
        pthreads_mutex_lock(lock);
        while(count == MAX_SIZE)
            pthread_cond_wait(notFull, lock);
        buf[count] = getChar();
        count++;
        pthread_cond_signal(notEmpty);
        pthread_mutex_unlock(lock);
    }
}

void consumer(char* buf) {
    for(;;) {
        pthread_mutex_lock(lock);
        while(count == 0)
            pthread_cond_wait(notEmpty, lock);
        useChar(buf[count-1]);
        count--;
        pthread_cond_signal(notFull);
        pthread_mutex_unlock(lock);
    }
}

int main() {
    char buffer[MAX_SIZE];
    pthread_t p;
    count = 0;
    pthread_mutex_init(&bufLock);
    pthread_cond_init(&notFull);
    pthread_cond_init(&notEmpty);
    pthread_create(&p, NULL, (void*)producer, &buffer);
    consume(&buffer);
    return 0;
}

```

31

More on Signaling Threads

- The previous example only wakes a single thread
 - not much control over which thread this is
- Perhaps all threads waiting on a condition need to be woken up
 - can do a broadcast of a signal
 - very similar to a regular signal in every other respect
- Prototype:
 - `int pthread_cond_broadcast(pthread_cond_t *cv);`
 - `cv`: condition variable to signal all waiters on

32